LOKESH D

Data Engineer | Databricks · PySpark · SQL · Azure · AWS · Data Warehousing

Chennai · 48lokesh@gmail.com · +91 99418 94248 · linkedin.com/in/lokesh-devaraj-data · github.com/Lokesh10Ki

PROFESSIONAL SUMMARY

Databricks Certified Data Engineer with 2.5+ years building ETL/ELT pipelines, Lakehouse architectures and data warehouse solutions using Apache Spark, PySpark, SQL and Databricks on Azure and AWS. Delivered production data platforms serving 17 business entities across 10+ source systems and currently migrating an enterprise Databricks Lakehouse from Azure to AWS with Unity Catalog governance. Strong in dimensional modelling, data quality, data lineage and CI/CD-driven deployments for financial and operational analytics.

TECHNICAL SKILLS

Data Engineering: Databricks, Apache Spark, PySpark, Spark SQL, Delta Lake, Unity Catalog, Databricks Workflows, ETL/ELT, dbt, Medallion Architecture, Lakehouse Architecture, Data Warehousing, Dimensional Modelling, CDC/SCD Type 2, Data Quality & Validation, Data Migration & Reconciliation

Big Data Ecosystem (Conceptual): Hadoop, HDFS, Hive, Apache Kafka, Apache Airflow

Cloud & Infrastructure: Azure (ADLS Gen2, Blob Storage, Key Vault, Azure DevOps), AWS (S3, IAM), Terraform, Databricks Asset Bundles, CI/CD, Docker

Programming: SQL (Advanced), Python (PySpark, Pandas, Automation Scripting), JavaScript, TypeScript

Business Intelligence: Power BI, DAX, Power Query (M), Row-Level Security (RLS), Semantic Models, Microsoft Fabric, Power BI Copilot, KPI Design, Power BI MCP

AI & Modern Data: Retrieval-Augmented Generation (RAG), Databricks Mosaic AI, Claude MCP, OpenAI API, Chatbot Integration

Development: React, Next.js, Node.js, NestJS, REST APIs, PostgreSQL, Git

PROFESSIONAL EXPERIENCE

Presidio | Data Engineer

Chennai, India · August 2026 - Present

Azure to AWS Databricks Migration | Databricks · Unity Catalog · PySpark · AWS S3 · IAM · Terraform · Databricks Asset Bundles

- **Platform Re-architecture:** Migrating an enterprise Databricks Lakehouse from Azure to AWS, re-platforming ADLS Gen2 storage to Amazon S3 and refactoring PySpark / Spark SQL notebooks, jobs, dbfs-mounts and orchestration into Databricks Workflows on the AWS workspace.
- **Unity Catalog Governance:** Migrated workspace objects to Unity Catalog, configuring catalogs, schemas, external locations and storage credentials backed by AWS IAM roles, with fine-grained access controls.
- **Validation & Cutover:** Built automated source-to-target reconciliation (row counts, schema checks and aggregate/checksum comparisons) and ran parallel Azure-AWS runs to confirm data parity ahead of production cutover.
- **IaC & CI/CD:** Provisioned workspace resources with Terraform and deployed jobs through Databricks Asset Bundles in CI/CD pipelines, enabling repeatable promotions across dev, test and prod environments.

JMAN Group (Solutions Enabler) | Data Engineer

Chennai, India · January 2024 - July 2026

Financial Analytics Platform - 17-entity professional services group | Databricks · PySpark · Spark SQL · Delta Lake · Azure ADLS Gen 2

- **Lakehouse Engineering:** Developed and optimized ETL/ELT data pipelines across a Medallion Lakehouse Architecture (Bronze, Silver, Gold), integrating data from 10+ source systems into a centralized data platform across 17 business entities.
- **Data Modelling:** Designed fact-dimension models for revenue, utilisation and contribution reporting, improving query performance by 5× while enhancing data quality through schema validation, deduplication and SCD Type 2 implementation.
- **Business Logic Ownership:** Owned the implementation and validation of financial business logic across Spark pipelines and Power BI semantic models, partnering with business stakeholders on KPI accuracy.

- **Revenue Reconciliation & PE Exit Support:** Led revenue reconciliation for a private equity exit by building fact models, validating financial KPIs and delivering trusted reporting for due diligence. Identified and resolved critical GRR/NRR discrepancies to improve reporting accuracy.
- **Power BI Delivery:** Built 15+ Power BI dashboards covering GRR/NRR, ARR Bridge, CLTV, employee utilisation, contribution and client performance, implementing Row-Level Security (RLS) for secure enterprise reporting.
- **Team Leadership:** Led a 3-member engineering team onboarding 2 new client entities into the live platform, reducing ramp-up time by 50% through structured documentation.
- **Data Lifecycle Management:** Implemented automated data validation, reconciliation and quality assurance checks across reporting pipelines, significantly reducing recurring data quality issues.

AI-Enhanced Analytics Initiative (POC) - Databricks Mosaic AI · RAG · Claude MCP · Power BI · Vector Store · LangChain

- Built an end-to-end Retrieval-Augmented Generation (RAG) solution using Databricks Mosaic AI by implementing document ingestion, intelligent chunking, embedding generation, Databricks Vector Search, semantic retrieval, and LLM grounding to enable accurate natural language querying over enterprise financial data.
- Developed a Power BI AI Review Agent using Claude MCP that automated semantic model validation, DAX development, review, metadata generation and documentation, reducing report review effort by ~70% while improving analytics quality and developer productivity.
- Evaluated and supported adoption of Microsoft Fabric Workspace and Power BI Copilot, enabling natural language DAX generation and reducing visualisation creation time by 60%.

Equiflow - Private Equity Data Platform - Full Stack | React · Next.js · Node.js · NestJS · Azure Blob Storage · OpenAI API · Docker

- Constructed a multi-tenant full-stack platform on Azure Blob Storage managing versioned financial documents with RBAC, audit logging and submission lifecycle state machine (Draft, Submitted, Approved).
- Developed NestJS REST APIs with PostgreSQL supporting concurrent users; implemented Microsoft SSO authentication, role-based access control and file versioning with conflict resolution.
- Engineered an AI-powered analytics feature that parsed uploaded financial spreadsheets to automatically generate interactive dashboards, chart configurations and natural language insights using OpenAI API integration.

E D U C A T I O N

B.Tech - Information Technology · Easwari Engineering College, Chennai · 2020 - 2024

CGPA: 9.14 / 10.00

K E Y A C H I E V E M E N T S

- Received a Certificate of Appreciation from the CTO in recognition of exceptional contributions and impact across strategic data engineering and analytics project initiatives.
- Recognised as the team's primary resource for business logic & data quality, resolving critical data quality and reporting issues.
- Trusted to collaborate directly with business, finance and executive stakeholders to translate complex reporting requirements into reliable data solutions.

C E R T I F I C A T I O N S

- **Databricks Certified Data Engineer Associate** - (Link)